

Decision Tree and Random Forest

Prof.Mingkui Tan

South China University of Technology
Southern Artificial Intelligence Laboratory(SAIL)

November 8, 2017



Contents

- 1 Decision Tree
 - Introduction
 - Decision Tree Example
 - Principled Criterion

- 2 Random Forest
 - Bagging
 - Random Forest

Decision Tree Example

Learning the Play Tennis Decision Tree

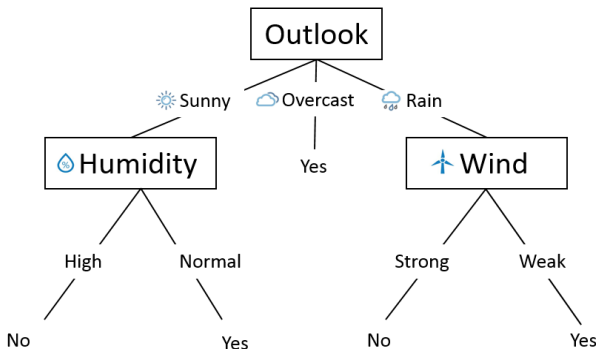
4 Attributes

Day	Outlook	Temp	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Decision Tree Example

Play Tennis: Yes or No?

- Learned function is a tree
- Classify instances by sorting them down from the root to the leaf node



Algorithm 1: Decision Tree

Input: Training set $D = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$; **Attribute set** $A = \{a_1, a_2, \dots, a_d\}$.

Procedure: Function TreeGenerate(D, A)

```

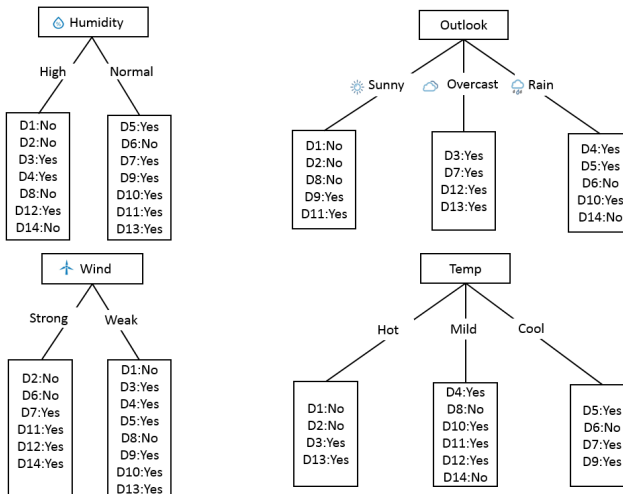
1  Generate node;
2  if all the samples in  $D$  belong to class  $C$  then
3  |   mark this node as class  $C$  leaf node; return
4  end
5  if  $A = \emptyset$  OR samples in  $D$  have the same value on  $A$  then
6  |   mark this node as leaf node, and the class should be the most frequent
   |   occurrence class; return
7  end
8  Select the best partition attribute  $a_*$  from  $A$ ;
9  for each value  $a_*^v$  of attribute  $a_*$  do
10 |    $D_v$  is the sample subset of  $D$  with  $a_* = a_*^v$ ;
11 |   if  $D_v = \emptyset$  then
12 |   |   mark this node as the leaf node, and the class should be the most frequent
   |   |   occurrence class; return
13 |   else
14 |   |   generate a branch for this node, TreeGenerate( $D_v, A \setminus \{a_*\}$ )
15 |   end
16 end

```

Output: A decision tree

Decision Tree Example

Which Is The Best Partition Attribute?



Decision Tree Example

A Good Attribute

An attribute is good when:

- For one value we get all instances as positive
- For other value we get all instances as negative

Decision Tree Example

Poor Attribute

An attribute is poor when:

- It provides no discrimination
- Attribute is immaterial to the decision
- For each value we have same number of positive and negative instances


Measure of Homogeneity of Examples

- **Entropy**: Characterizes the (im)purity of an arbitrary collection of examples.
- Given a collection D of positive and negative examples, entropy of D relative to boolean classification is

$$\text{Entropy}(D) \equiv -p_+ \log_2 p_+ - p_- \log_2 p_-$$

Where p_+ is proportion of positive examples and p_- is proportion of negative examples.

Entropy Function Relative to A Boolean Classification



Entropy

- Illustration:
- D is a collection of 14 examples with 9 positive and 5 negative examples
- Entropy of D relative to the Boolean classification:

$$Entropy(9+, 5-) = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940$$

- Entropy is zero if all members of D belong to the same class

Entropy for Multi-Valued Target function

- If the target attribute can take on c different values, the entropy of D relative to this c -wise classification is

$$Entropy(D) \equiv \sum_{i=1}^c -p_i \log_2 p_i$$

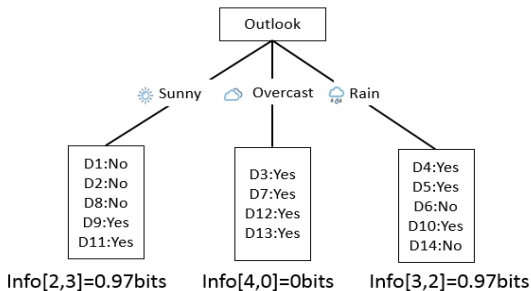
Information Gain Measures the Expected Reduction in Entropy

- Entropy measures the impurity of a collection
- Information gain of attribute A is the reduction in entropy caused by partitioning the set of examples D

$$Gain(D, A) \equiv Entropy(D) - \sum_{v \in Values(A)} \frac{|D_v|}{|D|} Entropy(D_v)$$

- Where $Values(A)$ is the set of all possible values for attribute A and D_v is the subset of D for which attribute A has value v

Measure of Purity: Information (bits)

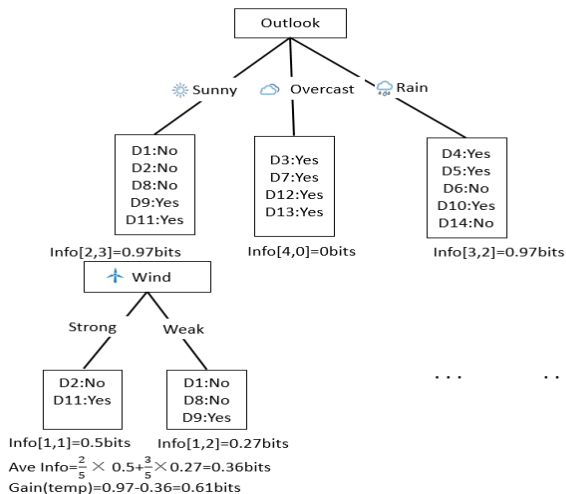


$$\begin{aligned}
 Info[2,3] &= entropy(2,3) \\
 &= -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} \\
 &= 0.97bits
 \end{aligned}$$

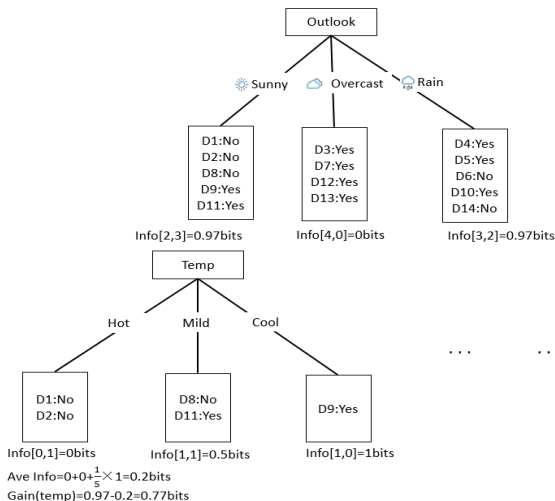
Information Gain for Each Attribute

- $\text{Gain}(\text{outlook}) = 0.94 - 0.693 = 0.247$
- $\text{Gain}(\text{temperature}) = 0.94 - 0.911 = 0.029$
- $\text{Gain}(\text{humidity}) = 0.94 - 0.788 = 0.152$
- $\text{Gain}(\text{windy}) = 0.94 - 0.892 = 0.048$
- $\arg \max_A \{0.247, 0.029, 0.152, 0.048\} = \text{outlook}$
- Select outlook as the splitting attribute of tree

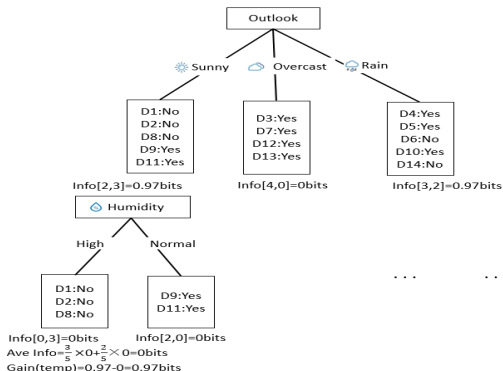
Expanded Tree Stumps for PlayTennis for Outlook=Sunny



Expanded Tree Stumps for PlayTennis for Outlook=Sunny

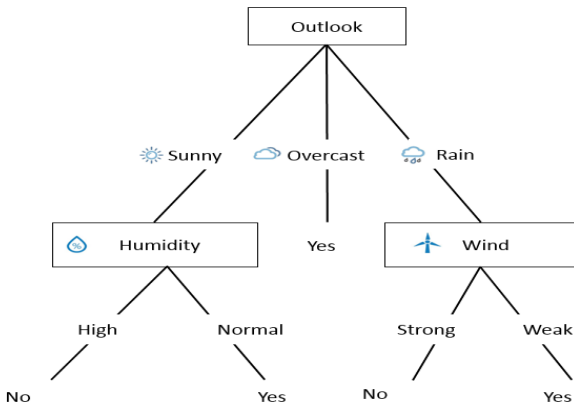


Expanded Tree Stumps for PlayTennis for Outlook=Sunny



- Since Gain(humidity) is highest, select humidity as splitting attribute
- No need to split further

Decision Tree for the Weather Data



Contents

- 1 Decision Tree
 - Introduction
 - Decision Tree Example
 - Principled Criterion
- 2 Random Forest
 - Bagging
 - Random Forest

Bagging

Bootstrap Sampling

- Giving a dataset D with m samples, we take samples to make dataset D'
- Select a sample from D to D' and put it back to the data set D each time, so the sample has the probability to be selected next time
- Repeat the procedure m times to get D' with m samples

Bagging

Bootstrap Sampling


The probability that a sample will not be selected:

$$\lim_{m \rightarrow +\infty} \left(1 - \frac{1}{m}\right)^m \rightarrow \frac{1}{e} \approx 0.368$$

- There are approximately $\frac{1}{3}$ samples used for testing which is called **out-of-bag estimate**
- Bootstrap sampling is useful in the case of small dataset and getting testing samples difficultly

Bagging

- Get T sampling sets through bootstrap sampling
- Train T base learners through the sampling sets respectively



- For classification:
The class with the most votes becomes the final class
- For regression:
The final output is the average output of every base learner

Random Forest

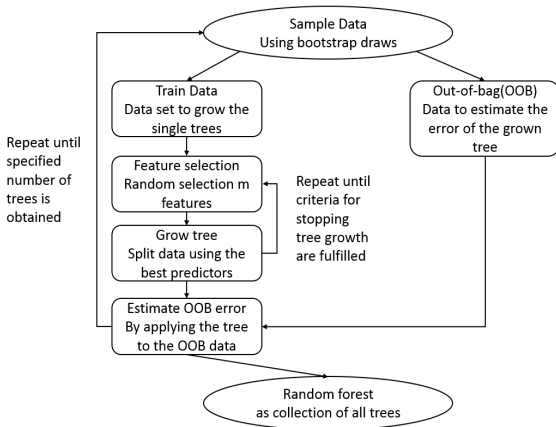
- Random forest is an extension of bagging using decision tree as base learner
- Randomly select **m out of p** features to get the optimal partition feature

Comparison between bagging and random forest

- The training efficiency of random forest is higher than bagging
- Bagging uses decision tree with **definite structure**
- Random forest uses decision tree with **random structure**

Random Forest

An example of the process flow is depicted below



THANK YOU!

Principal Component Analysis

Consider a data matrix, $\mathbf{X} \in \mathbb{R}^{n \times p}$, with **column-wise zero mean**, PCA is to find a direction $\mathbf{w} \in \mathbb{R}^p$ to project the data to the direction to achieve **the largest variance**

- Direction: $\|\mathbf{w}\| = 1$
- n observations: $y_i = \mathbf{x}^\top \mathbf{w}$, where $\mathbf{x} \in \mathbb{R}^p$ is a sample
- Projected feature: $\mathbf{y} = \mathbf{X}\mathbf{w} \in \mathbb{R}^n$
- \mathbf{w} is unknown
- Largest variance?

$$\max_{\mathbf{w}} \|\mathbf{X}\mathbf{w}\|^2 = \max_{\mathbf{w}} \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w}, s.t., \mathbf{w}^\top \mathbf{w} = 1 \quad (1)$$

Principal Component Analysis

$$\max_{\mathbf{w}} \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w}, s.t., \mathbf{w}^\top \mathbf{w} = 1 \quad (2)$$

- Lagrangian multiplier method:

$$\mathcal{L}(\mathbf{w}, \lambda) = \mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - \lambda(\mathbf{w}^\top \mathbf{w} - 1)$$

- Optimality condition:

$$\frac{\partial \mathcal{L}(\mathbf{w}, \lambda)}{\partial \mathbf{w}} = 0. \Rightarrow \mathbf{X}^\top \mathbf{X} \mathbf{w} - \lambda \mathbf{w} = \mathbf{0} \Rightarrow \mathbf{X}^\top \mathbf{X} \mathbf{w} = \lambda \mathbf{w}$$

Interpretations and Power Method

Rayleigh quotient:

$$\max_{\mathbf{w}} \frac{\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}}{\mathbf{w}^T \mathbf{w}} \quad (3)$$

- \mathbf{w} :
Direction or Principal Component or Eigenvector of $\mathbf{X}^T \mathbf{X}$.
- λ :
The objective value or Variance of y or Eigenvalue of $\mathbf{X}^T \mathbf{X}$.

Second Component?

- Compute the residual: $\mathbf{X}_r = \mathbf{X} - \mathbf{X}\mathbf{w}\mathbf{w}^\top$
- Rayleigh quotient on \mathbf{X}_r :

$$\max_{\mathbf{w}} \frac{\mathbf{w}^\top \mathbf{X}_r^\top \mathbf{X}_r \mathbf{w}}{\mathbf{w}^\top \mathbf{w}} \quad (4)$$

Extensions

Consider optimizing $\mathbf{W} \in \mathbb{R}^{p \times r}$ directly:

$$\max_{\mathbf{W} \in \mathbb{R}^{p \times r}} \|\mathbf{X}\mathbf{W}\|_F^2 = \mathbf{W}^\top \mathbf{X}^\top \mathbf{X} \mathbf{W}, s.t., \mathbf{W}^\top \mathbf{W} = \mathbf{I} \quad (5)$$

- Lagrangian multiplier method:
 $\mathcal{L}(\mathbf{W}, \lambda) = \mathbf{W}^\top \mathbf{X}^\top \mathbf{X} \mathbf{W} - \langle \Lambda, \mathbf{W}^\top \mathbf{W} - \mathbf{I} \rangle$
- Optimality condition: $\frac{\partial \mathcal{L}(\mathbf{W}, \lambda)}{\partial \mathbf{W}} = 0. \Rightarrow$
 $\mathbf{X}^\top \mathbf{X} \mathbf{W} - \lambda \mathbf{W} = 0 \Rightarrow \mathbf{X}^\top \mathbf{X} \mathbf{W} = \Lambda \mathbf{W}$
- If $r = p$, what will happen?